Informed search algorithms

Chapter 4, Sections 1–2, 4

Outline

- ♦ Best-first search
- \Diamond A* search
- ♦ Heuristics
- ♦ Hill-climbing
- ♦ Simulated annealing

Review: General search

```
 \begin{aligned} & \textbf{function} \; \text{General-Search}(\; problem, \text{Queuing-Fn}) \; \textbf{returns} \; \text{a solution, or failure} \\ & \textit{nodes} \leftarrow \text{Make-Queue}(\text{Make-Node}(\text{Initial-State}[problem])) \\ & \textbf{loop do} \\ & \quad \textbf{if} \; nodes \; \text{is empty then return} \; \text{failure} \\ & \textit{node} \leftarrow \text{Remove-Front}(\textit{nodes}) \\ & \quad \textbf{if} \; \text{Goal-Test}[problem] \; \text{applied to State}(\textit{node}) \; \text{succeeds then return} \; \textit{node} \\ & \textit{nodes} \leftarrow \text{Queuing-Fn}(\textit{nodes}, \text{Expand}(\textit{node}, \text{Operators}[problem])) \\ & \quad \textbf{end} \end{aligned}
```

A strategy is defined by picking the *order* of node expansion

Best-first search

Idea: use an evaluation function for each node

- estimate of "desirability"
- ⇒ Expand most desirable unexpanded node

Implementation:

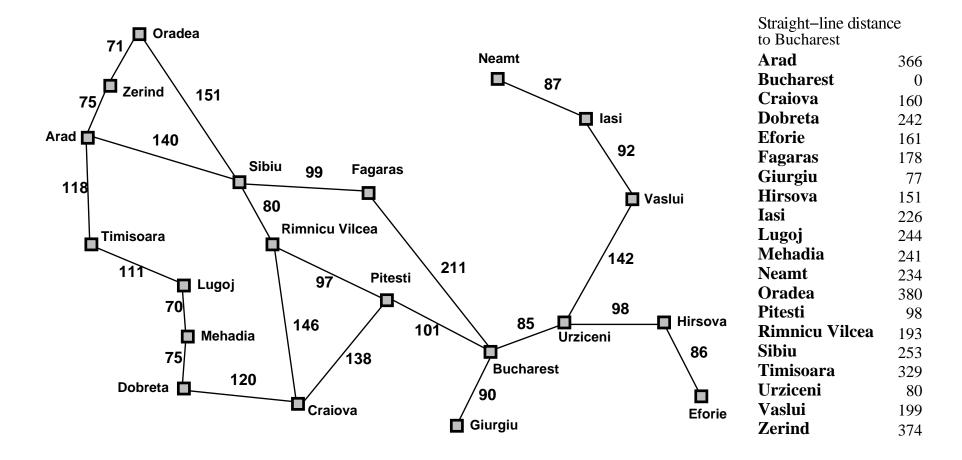
QUEUEINGFN = insert successors in decreasing order of desirability

Special cases:

greedy search

A* search

Romania with step costs in km



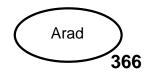
Greedy search

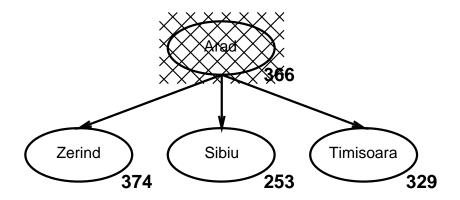
Evaluation function h(n) (heuristic) = estimate of cost from n to goal

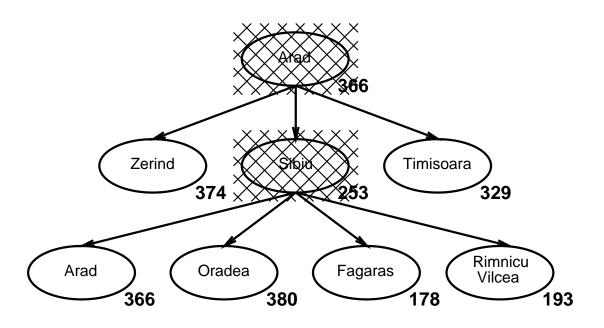
E.g., $h_{SLD}(n) = \text{straight-line distance from } n \text{ to Bucharest}$

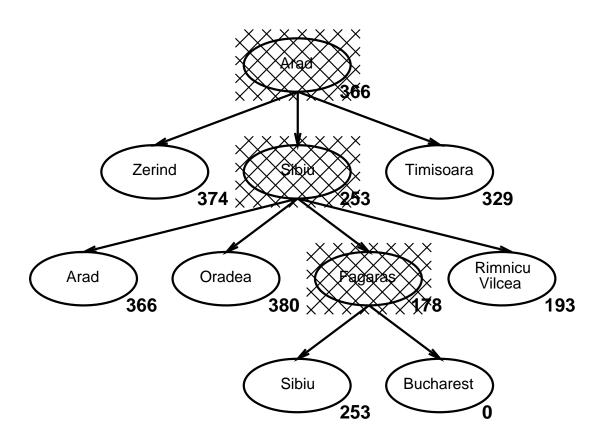
Greedy search expands the node that appears to be closest to goal

Greedy search example









Properties of greedy search

Complete??

Time??

Space??

Optimal??

Properties of greedy search

Complete?? No-can get stuck in loops, e.g.,

 $\mathsf{lasi} \to \mathsf{Neamt} \to \mathsf{lasi} \to \mathsf{Neamt} \to$

Complete in finite space with repeated-state checking

<u>Time</u>?? $O(b^m)$, but a good heuristic can give dramatic improvement

Space?? $O(b^m)$ —keeps all nodes in memory

Optimal?? No

A^* search

Idea: avoid expanding paths that are already expensive

Evaluation function f(n) = g(n) + h(n)

 $g(n) = \cos t$ so far to reach n

h(n) =estimated cost to goal from n

f(n) =estimated total cost of path through n to goal

A* search uses an admissible heuristic

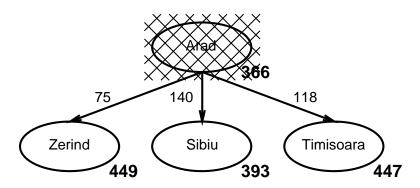
i.e., $h(n) \leq h^*(n)$ where $h^*(n)$ is the $true \cos t$ from n.

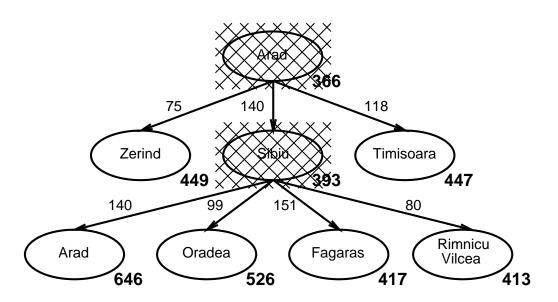
E.g., $h_{\rm SLD}(n)$ never overestimates the actual road distance

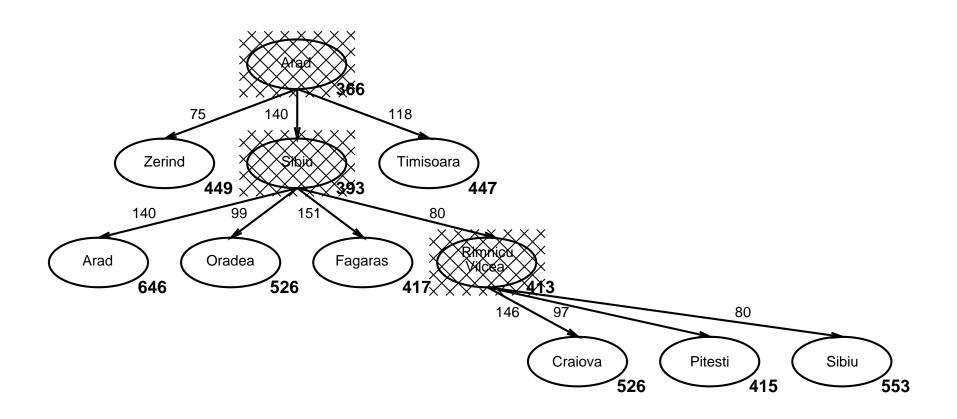
Theorem: A* search is optimal

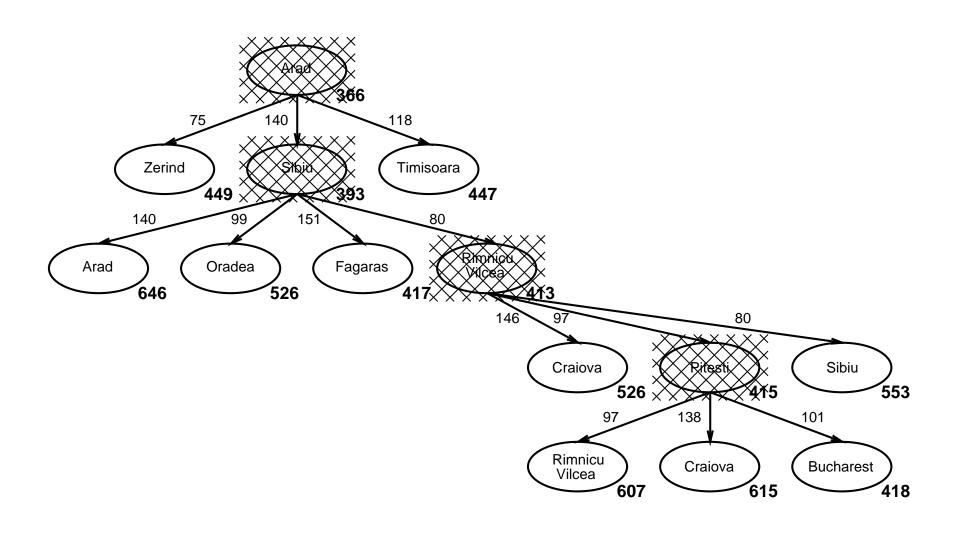
A^* search example

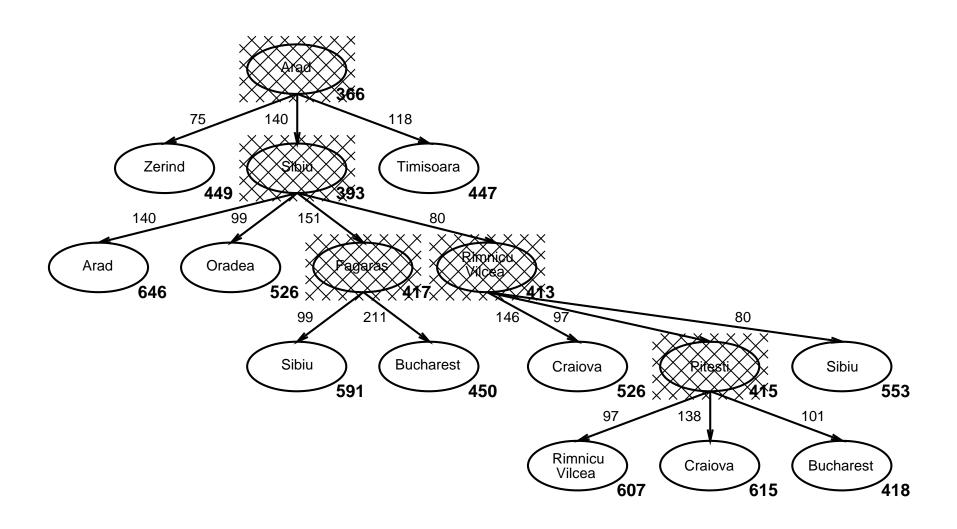






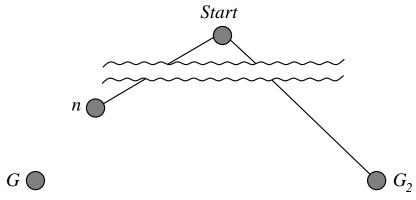






Optimality of A^* (standard proof)

Suppose some suboptimal goal G_2 has been generated and is in the queue. Let n be an unexpanded node on a shortest path to an optimal goal G_1 .



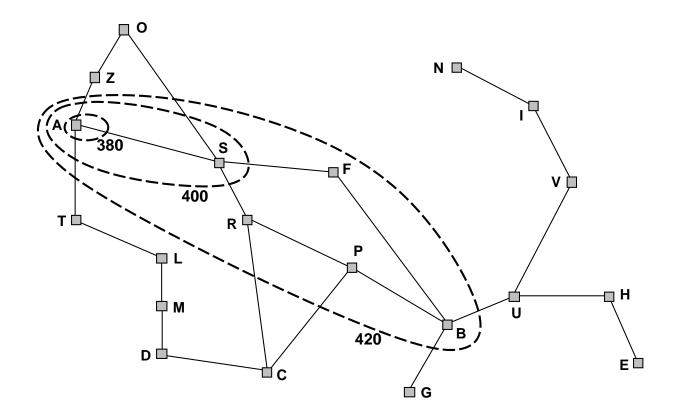
$$f(G_2) = g(G_2)$$
 since $h(G_2) = 0$
> $g(G_1)$ since G_2 is suboptimal
 $\geq f(n)$ since h is admissible

Since $f(G_2) > f(n)$, A* will never select G_2 for expansion

Optimality of A* (more useful)

<u>Lemma</u>: A^* expands nodes in order of increasing f value

Gradually adds "f-contours" of nodes (cf. breadth-first adds layers) Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$



Properties of A^*

<u>Complete</u>?? Yes, unless there are infinitely many nodes with $f \leq f(G)$

<u>Time</u>?? Exponential in [relative error in $h \times \text{length of soln.}$]

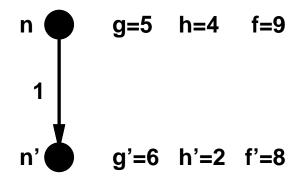
Space?? Keeps all nodes in memory

Optimal?? Yes—cannot expand f_{i+1} until f_i is finished

Proof of lemma: Pathmax

For some admissible heuristics, f may decrease along a path

E.g., suppose n' is a successor of n



But this throws away information! $f(n) = 9 \Rightarrow$ true cost of a path through n is ≥ 9 Hence true cost of a path through n' is ≥ 9 also

Pathmax modification to A*:

Instead of
$$f(n') = g(n') + h(n')$$
, use $f(n') = max(g(n') + h(n'), f(n))$

With pathmax, f is always nondecreasing along any path

Admissible heuristics

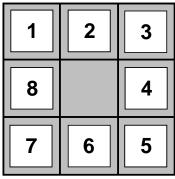
E.g., for the 8-puzzle:

 $h_1(n) = \text{number of misplaced tiles}$

 $h_2(n) = \text{total } \underline{\text{Manhattan}} \text{ distance}$

(i.e., no. of squares from desired location of each tile)

5	4	
6	1	8
7	3	2



Goal State

$$\underline{\frac{h_1(S) =??}{h_2(S) =??}}$$

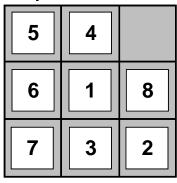
Admissible heuristics

E.g., for the 8-puzzle:

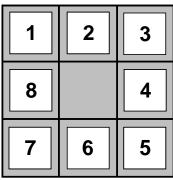
 $h_1(n) = \text{number of misplaced tiles}$

 $h_2(n) = \text{total } \underline{\text{Manhattan}} \text{ distance}$

(i.e., no. of squares from desired location of each tile)



Start State



Goal State

$$\underline{\frac{h_1(S)}{h_2(S)}} = ??? 7$$

 $\underline{h_2(S)} = ?? 2+3+3+2+4+2+0+2 = 18$

Dominance

If $h_2(n) \ge h_1(n)$ for all n (both admissible) then h_2 dominates h_1 and is better for search

Typical search costs:

$$d=14$$
 IDS = 3,473,941 nodes $A^*(h_1)=539$ nodes $A^*(h_2)=113$ nodes $d=14$ IDS = too many nodes $A^*(h_1)=39,135$ nodes $A^*(h_2)=1,641$ nodes

Relaxed problems

Admissible heuristics can be derived from the exact solution cost of a relaxed version of the problem

If the rules of the 8-puzzle are relaxed so that a tile can move anywhere, then $h_1(n)$ gives the shortest solution

If the rules are relaxed so that a tile can move to any adjacent square, then $h_2(n)$ gives the shortest solution

For TSP: let path be any structure that connects all cities \implies minimum spanning tree heuristic

Iterative improvement algorithms

In many optimization problems, path is irrelevant; the goal state itself is the solution

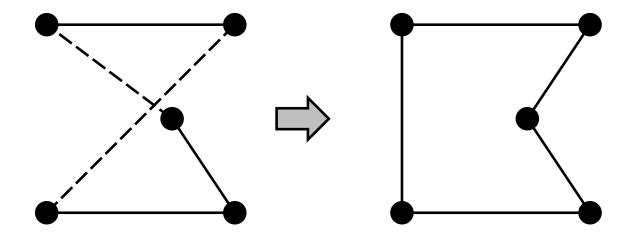
Then state space = set of "complete" configurations; find optimal configuration, e.g., TSP or, find configuration satisfying constraints, e.g., n-queens

In such cases, can use *iterative improvement* algorithms; keep a single "current" state, try to improve it

Constant space, suitable for online as well as offline search

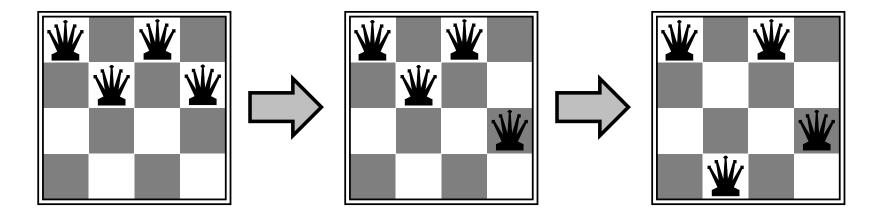
Example: Travelling Salesperson Problem

Find the shortest tour that visits each city exactly once



Example: *n*-queens

Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

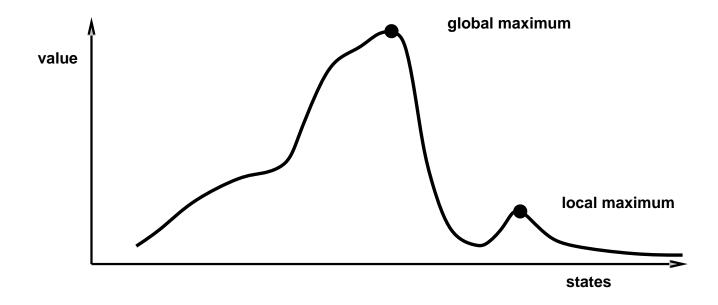


Hill-climbing (or gradient ascent/descent)

"Like climbing Everest in thick fog with amnesia"

Hill-climbing contd.

Problem: depending on initial state, can get stuck on local maxima



Simulated annealing

Idea: escape local maxima by allowing some "bad" moves but gradually decrease their size and frequency

```
function SIMULATED-ANNEALING( problem, schedule) returns a solution state inputs: problem, a problem schedule, a mapping from time to "temperature" local variables: current, a node next, a node T, a "temperature" controlling the probability of downward steps  \begin{array}{c} current \leftarrow \text{MAKE-NODE}(\text{INITIAL-STATE}[problem]) \\ \text{for } t\leftarrow 1 \text{ to} \propto \text{do} \\ T\leftarrow schedule[t] \\ \text{if } T{=}0 \text{ then return } current \\ next\leftarrow \text{a randomly selected successor of } current \\ \Delta E \leftarrow \text{VALUE}[next] - \text{VALUE}[current] \\ \text{if } \Delta E > 0 \text{ then } current \leftarrow next \\ \text{else } current \leftarrow next \text{ only with probability } e^{\Delta E/T} \\ \end{array}
```

Properties of simulated annealing

At fixed "temperature" T, state occupation probability reaches Boltzman distribution

$$p(x) = \alpha e^{\frac{E(x)}{kT}}$$

T decreased slowly enough \Longrightarrow always reach best state

Is this necessarily an interesting guarantee??

Devised by Metropolis et al., 1953, for physical process modelling

Widely used in VLSI layout, airline scheduling, etc.